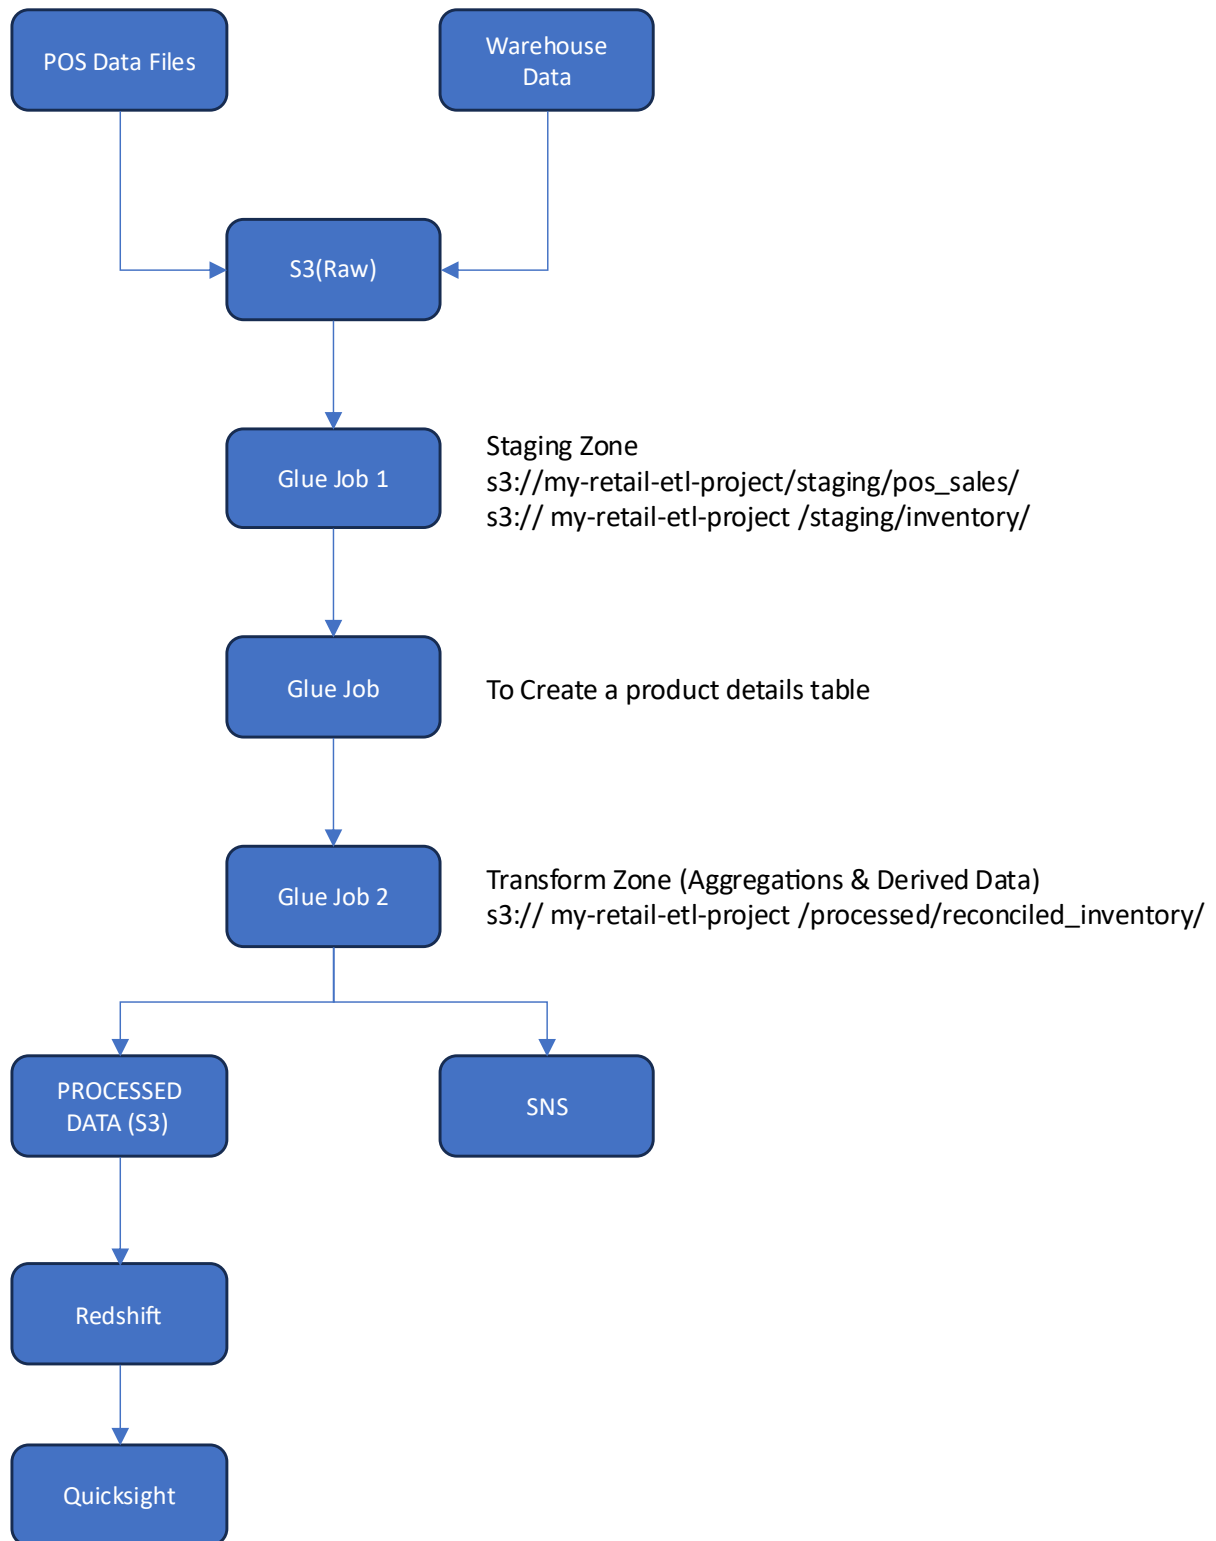


Retail Inventory Reconciliation (Batch ETL)

Summary:

- **Purpose:** This document outlines the architecture and data flow for the automated sales data pipeline. The primary goal of this pipeline is to ingest Point-of-Sale (POS) and supplementary warehouse data, process it, and load it into a centralized data warehouse to enable business intelligence (BI) and reporting.
- **Business Goal:** To provide the sales and marketing teams with daily updated dashboards in Amazon QuickSight for analyzing sales performance, product trends, and store efficiency.
- **Technologies Used:** AWS S3, AWS Glue, AWS SNS, Amazon Redshift, Amazon QuickSight, Airflow.

Architecture Diagram:

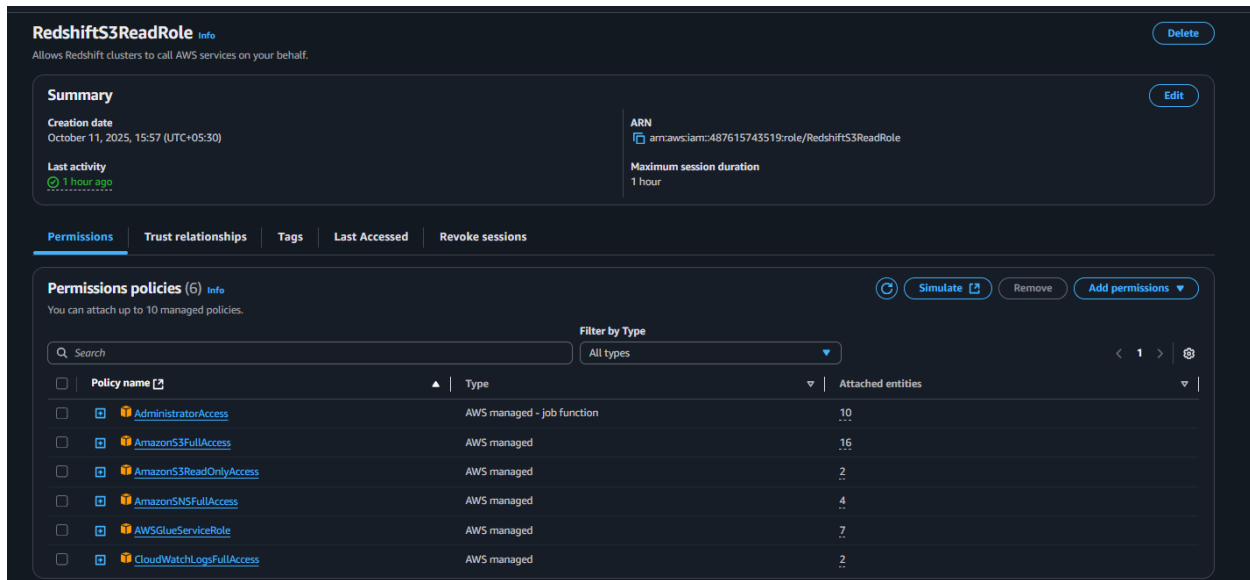


Detailed Implementation Steps & Configuration:

Step 1: Prerequisite: IAM Role Configuration

Before creating the pipeline, we need IAM roles with the correct permissions to allow AWS services to interact with each other.

IAM Role:



Step 2: Amazon S3 Bucket Setup

First, we'll create the main container for your project files. Remember, S3 bucket names must be **globally unique**.

1. **Navigate to S3:** In the AWS Management Console, search for and select "S3".
2. **Start Bucket Creation:** Click the **Create bucket** button.
3. **Configure Bucket Name and Region:**
 - **Bucket name:** Choose a unique name. A good practice is [project-name]-[purpose]-[date].
 - Example: retail-etl-project-datalake-20251013
 - **AWS Region:** Select the region closest to you, for example, **Asia Pacific (Mumbai) ap-south-1**.
4. **Object Ownership:** Leave this as the recommended default, **ACLs disabled**.

5. Block Public Access (Security):

- Keep the box for **Block all public access** checked. This is the most secure setting and is correct for a data pipeline where access is programmatic, not public.

6. Bucket Versioning (Data Protection):

- Select **Enable**. This is a best practice that keeps a history of all your file versions, protecting you from accidental overwrites or deletions.

7. Create the Bucket:

- Scroll to the bottom and click **Create bucket**.

You now have a secure, versioned S3 bucket ready for your data.

```
raw/
├── pos_sales/
│   └── date=2025-10-13/
│       └── pos_sales_2025-10-13.csv      <- Sales for Oct 13th
├── warehouse_inventory/
│   ├── date=2025-10-12/
│   │   └── warehouse_inventory_2025-10-12.csv  <- Opening Stock for Oct 13th
│   └── date=2025-10-13/
│       └── warehouse_inventory_2025-10-13.csv  <- Closing Stock for Oct 13th
```

General purpose buckets All AWS Regions **Directory buckets**

General purpose buckets (2) Info Copy ARN Empty Delete Create bucket

Buckets are containers for data stored in S3.

☐

[aws-glue-assets-487615743519-ap-south-1](#)

☐

[my-retail-etl-project](#)

AWS Region

Asia Pacific (Mumbai) ap-south-1

Creation date

October 13, 2025, 12:43:52 (UTC+05:30)

Asia Pacific (Mumbai) ap-south-1

October 13, 2025, 12:35:41 (UTC+05:30)

my-retail-etl-project Info

Objects Properties Permissions Metrics Management Access Points

Objects (4) Copy S3 URI Copy URL Download Open Delete Actions Create folder Upload

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 Inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

<input type="checkbox"/>	Name	Type	Last modified	Size	Storage class
<input type="checkbox"/>	processed_\$folder\$	-	October 13, 2025, 13:30:43 (UTC+05:30)	-	0 B Standard
<input type="checkbox"/>	processed/	Folder	-	-	-
<input type="checkbox"/>	raw/	Folder	-	-	-
<input type="checkbox"/>	staging/	Folder	-	-	-

raw/

Copy S3 URI

Objects

Properties

Objects (2)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Find objects by prefix

<input type="checkbox"/>	Name	Type	Last modified	Size	Storage class
<input type="checkbox"/>	pos_sales/	Folder	-	-	-
<input type="checkbox"/>	warehouse_inventory/	Folder	-	-	-

Amazon S3 > Buckets > my-retail-etl-project > raw/ > pos_sales/

pos_sales/

Copy S3 URI

Objects

Properties

Objects (1)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Find objects by prefix

<input type="checkbox"/>	Name	Type	Last modified	Size	Storage class
<input type="checkbox"/>	date=2025-10-13/	Folder	-	-	-

Amazon S3 > Buckets > my-retail-etl-project > raw/ > pos_sales/ > date=2025-10-13/

date=2025-10-13/

Copy S3 URI

Objects

Properties

Objects (1)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Find objects by prefix

<input type="checkbox"/>	Name	Type	Last modified	Size	Storage class
<input type="checkbox"/>	pos_sales_2025-10-13.csv	csv	October 13, 2025, 19:48:15 (UTC+05:30)	2.4 KB	Standard

Amazon S3 > Buckets > my-retail-etl-project > raw/ > warehouse_inventory/

warehouse_inventory/

Copy S3 URI

Objects

Properties

Objects (2)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Find objects by prefix

<input type="checkbox"/>	Name	Type	Last modified	Size	Storage class
<input type="checkbox"/>	date=2025-10-12/	Folder	-	-	-
<input type="checkbox"/>	date=2025-10-13/	Folder	-	-	-

Amazon S3 > Buckets > my-retail-etl-project > raw/ > warehouse_inventory/ > date=2025-10-12/

date=2025-10-12/

Copy S3 URI

Objects

Properties

Objects (1)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Find objects by prefix

<input type="checkbox"/>	Name	Type	Last modified	Size	Storage class
<input type="checkbox"/>	warehouse_inventory_2025-10-12.csv	csv	October 13, 2025, 19:48:54 (UTC+05:30)	2.7 KB	Standard

Amazon S3 > Buckets > my-retail-etl-project > raw/ > warehouse_inventory/ > date=2025-10-13/

date=2025-10-13/

Copy S3 URI

Objects

Properties

Objects (1)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Find objects by prefix

<input type="checkbox"/>	Name	Type	Last modified	Size	Storage class
<input type="checkbox"/>	warehouse_inventory_2025-10-13.csv	csv	October 13, 2025, 19:49:10 (UTC+05:30)	2.7 KB	Standard

pos_sales_2025-10-13:

```
transaction_id,store_id,sku,product_name,sale_date,quantity,price
T051,S001,SKU001,Wireless Mouse,2025-10-13,3,599
T052,S002,SKU008,Headphones,2025-10-13,5,1299
T053,S003,SKU016,Monitor 27Inch,2025-10-13,1,10999
T054,S004,SKU022,Microphone,2025-10-13,2,2999
T055,S001,SKU030,VR Headset,2025-10-13,1,9999
T056,S002,SKU045,Headphones,2025-10-13,4,1299
T057,S003,SKU012,Pen Drive 32GB,2025-10-13,8,499
T058,S004,SKU004,HDMI Cable,2025-10-13,12,299
T059,S001,SKU009,Mouse Pad,2025-10-13,10,199
T060,S002,SKU019,Portable SSD,2025-10-13,2,7999
T061,S003,SKU028,Stylus Pen,2025-10-13,5,599
T062,S004,SKU036,Wifi Adapter,2025-10-13,3,799
T063,S001,SKU050,Router,2025-10-13,1,2399
T064,S002,SKU041,HDMI Cable,2025-10-13,7,299
T065,S003,SKU033,Screen Cleaner Kit,2025-10-13,4,299
T066,S004,SKU025,Power Bank,2025-10-13,3,1299
T067,S001,SKU017,Smart Watch,2025-10-13,1,4999
T068,S002,SKU002,Keyboard,2025-10-13,4,899
T069,S003,SKU011,External HDD 1TB,2025-10-13,2,4599
T070,S004,SKU020,Wireless Charger,2025-10-13,6,999
T071,S001,SKU031,Smartphone Stand,2025-10-13,5,499
T072,S002,SKU048,External HDD 1TB,2025-10-13,1,4599
T073,S003,SKU039,Keyboard,2025-10-13,3,899
T074,S004,SKU049,Pen Drive 32GB,2025-10-13,10,499
T075,S001,SKU005,USB Hub,2025-10-13,2,499
T076,S002,SKU015,Laptop Bag,2025-10-13,2,1499
T077,S003,SKU024,Phone Case,2025-10-13,15,299
T078,S004,SKU032,Wireless Earbuds,2025-10-13,3,3499
T079,S001,SKU042,USB Hub,2025-10-13,4,499
T080,S002,SKU007,WebCam,2025-10-13,1,1999
T081,S003,SKU014,Printer Ink,2025-10-13,2,699
T082,S004,SKU021,HD Webcam,2025-10-13,2,2499
T083,S001,SKU035,Laptop Cooling Pad,2025-10-13,1,1499
T084,S002,SKU044,Webcam,2025-10-13,3,1999
T085,S003,SKU003,Monitor 24Inch,2025-10-13,1,7999
T086,S004,SKU013,Router,2025-10-13,1,2399
T087,S001,SKU023,Speaker Bluetooth,2025-10-13,2,1799
T088,S002,SKU034,Phone Tripod,2025-10-13,2,799
T089,S003,SKU043,Laptop Stand,2025-10-13,4,999
T090,S004,SKU047,Wireless Keyboard,2025-10-13,2,1299
T091,S001,SKU010,Wireless Keyboard,2025-10-13,3,1299
T092,S002,SKU026,Extension Board,2025-10-13,5,899
T093,S003,SKU037,Graphic Tablet,2025-10-13,1,8999
T094,S004,SKU046,Mouse Pad,2025-10-13,8,199
T095,S001,SKU006,Laptop Stand,2025-10-13,3,999
T096,S002,SKU018,Tablet 10Inch,2025-10-13,1,14999
T097,S003,SKU027,Monitor Stand,2025-10-13,1,1199
T098,S004,SKU029,HDMI Splitter,2025-10-13,2,699
T099,S001,SKU038,Wireless Mouse,2025-10-13,3,599
T100,S002,SKU040,Monitor 24Inch,2025-10-13,1,7999
```

warehouse_inventory_2025-10-12:

```
transaction_id,store_id,sku,product_name,sale_date,quantity,price
T051,S001,SKU001,Wireless Mouse,2025-10-13,3,599
T052,S002,SKU008,Headphones,2025-10-13,5,1299
T053,S003,SKU016,Monitor 27Inch,2025-10-13,1,10999
T054,S004,SKU022,Microphone,2025-10-13,2,2999
T055,S001,SKU030,VR Headset,2025-10-13,1,9999
T056,S002,SKU045,Headphones,2025-10-13,4,1299
T057,S003,SKU012,Pen Drive 32GB,2025-10-13,8,499
T058,S004,SKU004,HDMI Cable,2025-10-13,12,299
T059,S001,SKU009,Mouse Pad,2025-10-13,10,199
T060,S002,SKU019,Portable SSD,2025-10-13,2,7999
T061,S003,SKU028,Stylus Pen,2025-10-13,5,599
T062,S004,SKU036,Wifi Adapter,2025-10-13,3,799
T063,S001,SKU050,Router,2025-10-13,1,2399
T064,S002,SKU041,HDMI Cable,2025-10-13,7,299
T065,S003,SKU033,Screen Cleaner Kit,2025-10-13,4,299
T066,S004,SKU025,Power Bank,2025-10-13,3,1299
T067,S001,SKU017,Smart Watch,2025-10-13,1,4999
T068,S002,SKU002,Keyboard,2025-10-13,4,899
T069,S003,SKU011,External HDD 1TB,2025-10-13,2,4599
T070,S004,SKU020,Wireless Charger,2025-10-13,6,999
T071,S001,SKU031,Smartphone Stand,2025-10-13,5,499
T072,S002,SKU048,External HDD 1TB,2025-10-13,1,4599
T073,S003,SKU039,Keyboard,2025-10-13,3,899
T074,S004,SKU049,Pen Drive 32GB,2025-10-13,10,499
T075,S001,SKU005,USB Hub,2025-10-13,2,499
T076,S002,SKU015,Laptop Bag,2025-10-13,2,1499
T077,S003,SKU024,Phone Case,2025-10-13,15,299
T078,S004,SKU032,Wireless Earbuds,2025-10-13,3,3499
T079,S001,SKU042,USB Hub,2025-10-13,4,499
T080,S002,SKU007,WebCam,2025-10-13,1,1999
T081,S003,SKU014,Printer Ink,2025-10-13,2,699
T082,S004,SKU021,HD Webcam,2025-10-13,2,2499
T083,S001,SKU035,Laptop Cooling Pad,2025-10-13,1,1499
T084,S002,SKU044,Webcam,2025-10-13,3,1999
T085,S003,SKU003,Monitor 24Inch,2025-10-13,1,7999
T086,S004,SKU013,Router,2025-10-13,1,2399
T087,S001,SKU023,Speaker Bluetooth,2025-10-13,2,1799
T088,S002,SKU034,Phone Tripod,2025-10-13,2,799
T089,S003,SKU043,Laptop Stand,2025-10-13,4,999
T090,S004,SKU047,Wireless Keyboard,2025-10-13,2,1299
T091,S001,SKU010,Wireless Keyboard,2025-10-13,3,1299
T092,S002,SKU026,Extension Board,2025-10-13,5,899
T093,S003,SKU037,Graphic Tablet,2025-10-13,1,8999
T094,S004,SKU046,Mouse Pad,2025-10-13,8,199
T095,S001,SKU006,Laptop Stand,2025-10-13,3,999
T096,S002,SKU018,Tablet 10Inch,2025-10-13,1,14999
T097,S003,SKU027,Monitor Stand,2025-10-13,1,1199
T098,S004,SKU029,HDMI Splitter,2025-10-13,2,699
T099,S001,SKU038,Wireless Mouse,2025-10-13,3,599
T100,S002,SKU040,Monitor 24Inch,2025-10-13,1,7999
```

warehouse_inventory_2025-10-13:

```
record_id,sku,product_name,category,warehouse_location,stock_on_hand,received_date
R051,SKU001,Wireless Mouse,Accessories,WH1,111,2025-10-13
R052,SKU002,Keyboard,Accessories,WH1,73,2025-10-13
R053,SKU003,Monitor 24Inch,Displays,WH2,42,2025-10-13
R054,SKU004,HDMI Cable,Cables,WH3,278,2025-10-13
R055,SKU005,USB Hub,Accessories,WH2,134,2025-10-13
R056,SKU006,Laptop Stand,Accessories,WH1,101,2025-10-13
R057,SKU007,WebCam,Electronics,WH4,47,2025-10-13
R058,SKU008,Headphones,Audio,WH2,185,2025-10-13
R059,SKU009,Mouse Pad,Accessories,WH3,158,2025-10-13
R060,SKU010,Wireless Keyboard,Accessories,WH1,133,2025-10-13
R061,SKU011,External HDD 1TB,Storage,WH2,65,2025-10-13
R062,SKU012,Pen Drive 32GB,Storage,WH3,232,2025-10-13
R063,SKU013,Router,Networking,WH4,57,2025-10-13
R064,SKU014,Printer Ink,Printers,WH1,33,2025-10-13
R065,SKU015,Laptop Bag,Accessories,WH2,94,2025-10-13
R066,SKU016,Monitor 27Inch,Displays,WH3,27,2025-10-13
R067,SKU017,Smart Watch,Wearables,WH4,22,2025-10-13
R068,SKU018,Tablet 10Inch,Tablets,WH1,18,2025-10-13
R069,SKU019,Portable SSD,Storage,WH2,35,2025-10-13
R070,SKU020,Wireless Charger,Accessories,WH3,66,2025-10-13
R071,SKU021,HD Webcam,Electronics,WH1,40,2025-10-13
R072,SKU022,Microphone,Audio,WH4,31,2025-10-13
R073,SKU023,Speaker Bluetooth,Audio,WH2,104,2025-10-13
R074,SKU024,Phone Case,Accessories,WH3,127,2025-10-13
R075,SKU025,Power Bank,Electronics,WH1,71,2025-10-13
R076,SKU026,Extension Board,Accessories,WH2,60,2025-10-13
R077,SKU027,Monitor Stand,Displays,WH4,22,2025-10-13
R078,SKU028,Stylus Pen,Accessories,WH1,82,2025-10-13
R079,SKU029,HDMI Splitter,Cables,WH3,37,2025-10-13
R080,SKU030,VR Headset,Electronics,WH2,13,2025-10-13
R081,SKU031,Smartphone Stand,Accessories,WH1,118,2025-10-13
R082,SKU032,Wireless Earbuds,Audio,WH2,52,2025-10-13
R083,SKU033,Screen Cleaner Kit,Accessories,WH3,190,2025-10-13
R084,SKU034,Phone Tripod,Accessories,WH4,65,2025-10-13
R085,SKU035,Laptop Cooling Pad,Accessories,WH1,57,2025-10-13
R086,SKU036,Wifi Adapter,Networking,WH2,39,2025-10-13
R087,SKU037,Graphic Tablet,Electronics,WH3,18,2025-10-13
R088,SKU038,Wireless Mouse,Accessories,WH4,93,2025-10-13
R089,SKU039,Keyboard,Accessories,WH1,77,2025-10-13
R090,SKU040,Monitor 24Inch,Displays,WH2,48,2025-10-13
R091,SKU041,HDMI Cable,Cables,WH3,294,2025-10-13
R092,SKU042,USB Hub,Accessories,WH4,71,2025-10-13
R093,SKU043,Laptop Stand,Accessories,WH1,89,2025-10-13
R094,SKU044,Webcam,Electronics,WH2,50,2025-10-13
R095,SKU045,Headphones,Audio,WH3,162,2025-10-13
R096,SKU046,Mouse Pad,Accessories,WH4,147,2025-10-13
R097,SKU047,Wireless Keyboard,Accessories,WH1,69,2025-10-13
R098,SKU048,External HDD 1TB,Storage,WH2,57,2025-10-13
R099,SKU049,Pen Drive 32GB,Storage,WH3,212,2025-10-13
R100,SKU050,Router,Networking,WH4,52,2025-10-13
```


Step 3: Glue Job 1

```
import sys
from awsglue.utils import getResolvedOptions
from pyspark.context import SparkContext
from awsglue.context import GlueContext
from awsglue.job import Job
import pyspark.sql.functions as F

# --- Initialization ---
sc = SparkContext()
glueContext = GlueContext(sc)
spark = glueContext.spark_session
job = Job(glueContext)

# --- Parameters (Dynamic) ---
args = getResolvedOptions(sys.argv, ["JOB_NAME", "processing_date"])
job.init(args["JOB_NAME"], args)
processing_date = args["processing_date"]
s3_bucket = "my-retail-etl-project"

# --- Define Paths ---
input_path = f"s3://{s3_bucket}/raw/pos_sales/date={processing_date}/"
output_path = f"s3://{s3_bucket}/staging/pos_sales/date={processing_date}/"

# --- Read Raw Data ---
# Reading the CSV with the final, correct schema
raw_pos_df = spark.read.format("csv") \
    .option("header", "true").option("inferSchema", "true") \
    .load(input_path)

# --- Transformations ---
# --- MODIFIED: Use the new, correct source column names: 'sku' and 'quantity' ---
cleaned_df = raw_pos_df \
    .withColumn("sku", F.upper(F.trim(F.col("sku")))) \
    .withColumn("quantity", F.col("quantity").cast("int"))

# --- MODIFIED: Group by 'sku'. We are no longer renaming this column. ---
# We also drop product_name from the grouping as it is descriptive data.
aggregated_sales_df = cleaned_df.groupBy("sku") \
    .agg(F.sum("quantity").alias("total_quantity_sold")) \
    .withColumn("date", F.lit(processing_date))

# --- Final Selection ---
# The 'sku' column name already matches our Redshift table, so no rename is needed.
final_df = aggregated_sales_df.select(
    F.col("date").cast("date").alias("date_key"),
    F.col("sku"),
    F.col("total_quantity_sold")
)

# --- Write to Staging Zone ---
final_df.write \
    .mode("overwrite") \
    .parquet(output_path)

job.commit()
```

Your jobs (3) Info

Filter jobs by property

<input type="checkbox"/>	Job name	Type	Created by	Last modified	AWS Glue version
<input type="checkbox"/>	glue-job-2-reconcile-inventory	Glue ETL	Script	10/13/2025, 6:30:42 PM	5.0
<input type="checkbox"/>	glue-job-create-dim-products	Glue ETL	Script	10/13/2025, 5:54:37 PM	5.0
<input type="checkbox"/>	glue-job-1-stage-pos-sales	Glue ETL	Script	10/13/2025, 5:52:03 PM	5.0

glue-job-1-stage-pos-sales

Last modified on 10/13/2025, 7:50:40 PM

Actions Save Run

Script

Job details

Runs

Data quality

Schedules

Version Control

Script Info

```
1 import sys
2 from aws glue.utils import getResolvedOptions
3 from pyspark.context import SparkContext
4 from aws glue.context import GlueContext
5 from aws glue.job import Job
6 import pyspark.sql.functions as F
7
8 # --- Initialization ---
9 sc = SparkContext()
10 glueContext = GlueContext(sc)
11 spark = glueContext.spark_session
12 job = Job(glueContext)
13
14 # --- Parameters (Dynamic) ---
15 args = getResolvedOptions(sys.argv, ["JOB_NAME", "processing_date"])
16 job.init(args["JOB_NAME"], args)
17 processing_date = args["processing_date"]
18 s3_bucket = "my-retail-etl-project"
19
20 # --- Define Paths ---
21 input_path = f"s3://{s3_bucket}/raw/pos_sales/date={processing_date}/"
22 output_path = f"s3://{s3_bucket}/staging/pos_sales/date={processing_date}/"
23
24 # --- Read Raw Data ---
25 # Reading the CSV with the final, correct schema
26 raw_pos_df = spark.read.format("csv") \
27     .option("header", "true").option("inferSchema", "true") \
28     .load(input_path)
29
30 # --- Transformations ---
31 # --- MODIFIED: use the new, correct source column names: 'sku' and 'quantity' ---
32 --
```

Python Ln 1, Col 1 0 Errors 0 0 Warnings 0

glue-job-1-stage-pos-sales

Last modified on 10/13/2025, 7:50:40 PM

Actions Save Run

Script

Job details

Runs

Data quality

Schedules

Version Control

Basic properties Info

Name

glue-job-1-stage-pos-sales

Description - optional

Descriptions can be up to 2048 characters long.

IAM Role

Role assumed by the job with permission to access your data stores. Ensure that this role has permission to your Amazon S3 sources, targets, temporary directory, scripts, and any libraries used by the job.

glue-retail-etl-role

Type

The type of ETL job. This is set automatically based on the types of data sources you have selected.

Spark

Glue version Info

Glue 5.0 - Supports spark 3.5, Scala 2, Python 3

Language

Python 3

Worker type

Set the types of predefined worker that is allowed when a job runs.

G 1X
(4vCPU and 16GB RAM)

Automatically scale the number of workers

☐ AWS Glue will optimize costs and resource usage by dynamically scaling the number of workers up and down throughout the job run. Requires Glue 5.0 or later.

Step 4: Glue Job For dim-products table

glue-job-create-dim-products

```
import sys
from awsglue.utils import getResolvedOptions
from pyspark.context import SparkContext
from awsglue.context import GlueContext
from awsglue.job import Job
import pyspark.sql.functions as F

# --- Initialization ---
sc = SparkContext()
glueContext = GlueContext(sc)
spark = glueContext.spark_session
job = Job(glueContext)

# --- Parameters (Dynamic) ---
args = getResolvedOptions(sys.argv, ["JOB_NAME", "processing_date"])
job.init(args["JOB_NAME"], args)
processing_date = args["processing_date"]
s3_bucket = "my-retail-etl-project"

# --- Define Paths ---
input_path = f"s3://{s3_bucket}/raw/warehouse_inventory/date={processing_date}/"
output_path = f"s3://{s3_bucket}/processed/dim_products/"

# --- Read Raw Data ---
inventory_df = spark.read.format("csv") \
    .option("header", "true").option("inferSchema", "true") \
    .load(input_path)

# --- Select and Clean Dimension Columns ---
# MODIFIED: Correctly reads 'sku' from the source file. No rename is needed.
dim_products_df = inventory_df \
    .withColumn("sku", F.upper(F.trim(F.col("sku")))) \
    .withColumn("product_name", F.trim(F.col("product_name"))) \
    .withColumn("category", F.trim(F.col("category"))) \
    .select("sku", "product_name", "category") \
    .dropDuplicates(["sku"])

# --- Write to Processed Zone ---
dim_products_df.coalesce(1).write \
    .mode("overwrite") \
    .parquet(output_path)

job.commit()
```

glue-job-create-dim-products

Script

Job details

Runs

Data quality

Schedules

Version Control

Basic properties

Info

Name

glue-job-create-dim-products

Description - optional

Descriptions can be up to 2048 characters long.

IAM Role

Role assumed by the job with permission to access your data stores. Ensure that this role has permission to your Amazon S3 sources, targets, temporary directory, scripts, and any libraries used by the job.

glue-retail-etl-role

Type

The type of ETL job. This is set automatically based on the types of data sources you have selected.

Spark

Glue version

Info

Glue 5.0 - Supports spark 3.5, Scala 2, Python 3

Language

Python 3

Worker type

Set the type of predefined worker that is allowed when a job runs.

G 1X
(4vCPU and 16GB RAM)

Automatically scale the number of workers

☐ AWS Glue will optimize costs and resource usage by dynamically scaling the number of workers up and down throughout the job run. Requires Glue 3.0 or later.

Step 5: Glue Job 2

```
import sys
from awsglue.utils import getResolvedOptions
from pyspark.context import SparkContext
from awsglue.context import GlueContext
from awsglue.job import Job
import pyspark.sql.functions as F
from datetime import datetime, timedelta
import boto3

# --- Initialization ---
sc = SparkContext()
glueContext = GlueContext(sc)
spark = glueContext.spark_session
job = Job(glueContext)

# --- Parameters (Dynamic) ---
args = getResolvedOptions(sys.argv, [
    "JOB_NAME",
    "processing_date",
    "s3_bucket",
    "aws_region",
    "sns_topic_arn"
])
job.init(args["JOB_NAME"], args)

processing_date_str = args["processing_date"]
s3_bucket = args["s3_bucket"]
aws_region = args["aws_region"]
sns_topic_arn = args["sns_topic_arn"]

# --- Date Calculations ---
processing_date = datetime.strptime(processing_date_str, '%Y-%m-%d')
previous_date_str = (processing_date - timedelta(days=1)).strftime('%Y-%m-%d')

# --- S3 Paths ---
staging_sales_path = f"s3://{s3_bucket}/staging/pos_sales/date={processing_date_str}/"
raw_inventory_path_today = f"s3://{s3_bucket}/raw/warehouse_inventory/date={processing_date_str}/"
raw_inventory_path_yesterday = f"s3://{s3_bucket}/raw/warehouse_inventory/date={previous_date_str}/"
dim_products_path = f"s3://{s3_bucket}/processed/dim_products/"
processed_output_path =
f"s3://{s3_bucket}/processed/reconciled_inventory/date={processing_date_str}/"

# --- 1. Load DataFrames ---
daily_sales_df = spark.read.parquet(staging_sales_path)
opening_stock_df = spark.read.format("csv").option("header", "true").option("inferSchema",
"true").load(raw_inventory_path_yesterday).select(F.upper(F.trim(F.col("sku")))).alias("sku"),F.col("
stock_on_hand").alias("opening_stock"))
actual_closing_stock_df = spark.read.format("csv").option("header", "true").option("inferSchema",
"true").load(raw_inventory_path_today).select(F.upper(F.trim(F.col("sku")))).alias("sku"),F.col("stoc
k_on_hand").alias("actual_closing_stock"))
dim_products_df = spark.read.parquet(dim_products_path)

# --- 2. Join and Reconcile ---
inventory_df = opening_stock_df.join(actual_closing_stock_df, "sku", "full_outer")
reconciliation_df = inventory_df.join(daily_sales_df, "sku", "left")
reconciliation_with_names_df = reconciliation_df.join(dim_products_df.select("sku", "product_name"),
"sku", "left")

# --- 3. Calculate and Finalize ---
final_df = reconciliation_with_names_df.fillna(0, subset=['opening_stock', 'actual_closing_stock',
'total_quantity_sold']).withColumn("expected_closing_stock", (F.col("opening_stock") -
```

```

F.col("total_quantity_sold"))).withColumn("discrepancy", (F.col("actual_closing_stock") -
F.col("expected_closing_stock"))).withColumn("date", F.lit(processing_date_str))

# --- 4. Select and Rename Final Columns ---
final_df_selected =
final_df.select(F.col("date").cast("date").alias("date_key"),F.col("sku"),F.col("product_name"),F.co
l("opening_stock"),F.col("total_quantity_sold").alias("quantity_sold"),F.col("expected_closing_stock
"),F.col("actual_closing_stock"),F.col("discrepancy").alias("discrepancy_amount"))

# --- 5. Write to Processed Zone ---
final_df_selected.write.mode("overwrite").parquet(processed_output_path)

# --- 6. Check for Discrepancies and Send Notification ---
discrepancy_df = final_df_selected.filter(F.col("discrepancy_amount") != 0)
if discrepancy_df.count() > 0:
    print("Discrepancies found. Preparing to send SNS notification.")
    sns_client = boto3.client('sns', region_name=aws_region)
    discrepancy_examples = discrepancy_df.limit(5).collect()
    message = f"Inventory reconciliation for date {processing_date_str} found discrepancies.\n\n"
    message += f"Total items with discrepancies: {discrepancy_df.count()}\n\n"
    message += "Example Discrepancies:\n"
    for row in discrepancy_examples:
        product_name = row['product_name'] if row['product_name'] else "N/A"
        message += f"- SKU: {row['sku']], Product: {product_name}, Discrepancy:
{row['discrepancy_amount']}\n"
    message += f"\nFull report available at: {processed_output_path}"
    subject = f"Alert: Inventory Discrepancy Found for {processing_date_str}"
    response = sns_client.publish(TopicArn=sns_topic_arn, Message=message, Subject=subject)
    print(f"SNS notification sent! Message ID: {response['MessageId']}")
else:
    print("No discrepancies found. No notification sent.")

job.commit()

```

glue-job-2-reconcile-inventory

Last modified on 10/15/2025, 8:06:36 PM

Actions

Save

Run

Script

Job details

Runs

Data quality

Schedules

Version Control

Script

Info

```

10 # --- Initialization ---
11 sc = SparkContext()
12 glueContext = GlueContext(sc)
13 spark = glueContext.spark_session
14 job = Job(glueContext)
15
16 # --- Parameters (Dynamic) ---
17 args = getResolvedOptions(sys.argv, [
18     "JOB_NAME",
19     "processing_date",
20     "s3_bucket",
21     "aws_region",
22     "sns_topic_arn"
23 ])
24 job.init(args["JOB_NAME"], args)
25
26 processing_date_str = args["processing_date"]
27 s3_bucket = args["s3_bucket"]
28 aws_region = args["aws_region"]
29 sns_topic_arn = args["sns_topic_arn"]
30
31 # --- Date Calculations ---
32 processing_date = datetime.strptime(processing_date_str, '%Y-%m-%d')
33 previous_date_str = (processing_date - timedelta(days=1)).strftime('%Y-%m-%d')
34
35 # --- S3 Paths ---
36 staging_sales_path = f"s3://{s3_bucket}/staging/pos_sales/date={processing_date_str}/"

```

Python

Ln 28, Col 28

Errors: 0

Warnings: 0

Step 6: Create a Redshift Serverless

The first screenshot shows the AWS Redshift Serverless 'Serverless dashboard'. The dashboard provides an overview of namespaces and workgroups. Key metrics include:

- Namespace overview:** Total snapshots: 0, Databases in my account: 0, Databases requiring authorization: 0, Databases from other accounts: 0, Databases requiring association: 0.
- Namespaces / Workgroups:** A table showing the 'default-namespace' and 'default-workgroup', both with a status of 'Available'. It also includes line graphs for average query duration and average number of queries run daily.
- Queries metrics:** A section for monitoring queries, with tabs for 'Running and queued queries', 'Completed and failed queries', and 'Number of queries'.
- Total compute usage:** A section for visualizing costs, showing 'Total consumed RPU hours' and a message about retrieving total compute usage.
- Free trial:** Information about the free trial, including 'Free trial credits remaining' (\$289.60 out of \$500.00) and 'Free trial expiration' (December 06, 2025).

The second screenshot shows the 'default-namespace' configuration page. It provides detailed information about the namespace and its associated workgroup:

- General information:** Namespace: default-namespace, Namespace ID: 19a62f6b-276a-462e-954f-ec298a0b0bbd, Namespace ARN: arn:aws:redshift-serverless:ap-south-1:487615745519:namespace/19a62f6b-276a-462e-954f-ec298a0b0bbd, Namespace register status: Deregistered.
- Status:** Available.
- Admin user name:** admin.
- Database name:** dev.
- Total table count:** 4.
- Storage used:** 1.7 GB.
- Workgroup name:** default-workgroup, Status: Available.

Workgroup
Data backup
Database
Security and encryption
Datashares
Zero-ETL integrations
Resource policy
Tags

Permissions info

Manage IAM roles

IAM roles	Status	Amazon Resource Name (ARN)	Role type
RedshiftS3ReadRole	In-sync	arn:aws:iam::487615743519:role/RedshiftS3ReadRole	-
AmazonRedshift-CommandsAccessRole-20250929172455	In-sync	arn:aws:iam::487615743519:role/service-role/AmazonRedshift-CommandsAccessRole-20250929172455	Default

default-workgroup info

Actions

Query data

General information

Workgroup

default-workgroup

Namespace

default-namespace

Workgroup ARN

arn:aws:redshift-serverless:ap-south-1:487615743519:workgroup/7d151f7e-defe-4dd1-982c-57e7b1447886

Workgroup version

1.0.129791

Status

Available

Date created

October 13, 2025, 13:40 (UTC+05:30)

Base capacity

4 RPU

Custom domain name

-

Patch version

Patch 194

Track

new

Current

Endpoint

default-workgroup.487615743519.ap-south-1.redshift-serverless.amazonaws.com:5439/dev

JDBC URL

jdbc:redshift://default-workgroup.487615743519.ap-south-1.redshift-serverless.amazonaws.com:5439/dev

ODBC URL

Driver={Amazon Redshift (64)}; Server=default-workgroup.487615743519.ap-south-1.redshift-serverless.amazonaws.com; Database=dev

Publicly accessible turned on.

Publicly accessible for this Redshift instance turned on. This means this instance is now accessible to the public internet. Any external sources will be able to connect and access the data within this Redshift cluster. [Learn more](#)

Query and database monitoring

Data access

Limits

Performance

Tags

Network and security info

Edit

Virtual private cloud (VPC)

vpc-08960c33efb0c5db9

VPC endpoint ID

vpc-09b9a31669b321b1

VPC security group

sg-04699b8435729b2af

Subnet

subnet-04ef7ff1da40da7, subnet-0f414f41149d21d8, subnet-08b910b896c6f8e7e

SSL

Enabled

Enhanced VPC routing

Off

Publicly accessible

Turned On

IP address type

IPv4

Step 7: Create Tables in Redshift

Redshift query editor v2

Create

Load data

Filter resources

Serverless: default-workgroup

native databases (2)

dev

public

Tables

dim_products

fact_daily_sales

fact_inventory_reconciliation

Views

Functions

Stored procedures

sample_data_dev

external databases (1)

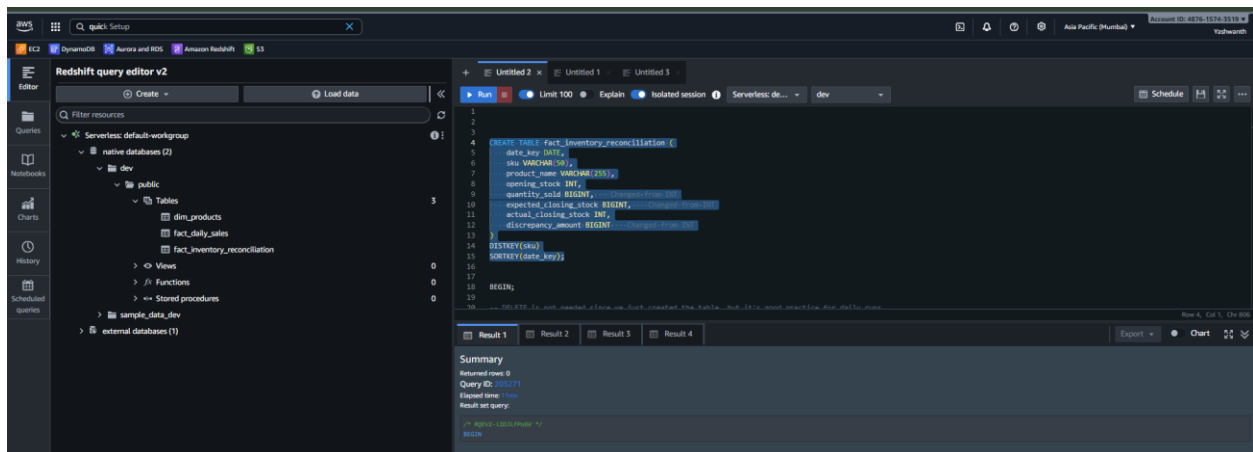

```

CREATE TABLE fact_inventory_reconciliation (
  date_key DATE,
  sku VARCHAR(50),
  product_name VARCHAR(255),
  opening_stock INT,
  quantity_sold BIGINT, -- Changed from INT
  expected_closing_stock BIGINT, -- Changed from INT
  actual_closing_stock INT,
  discrepancy_amount BIGINT -- Changed from INT
)
DISTKEY(sku)
SORTKEY(date_key);

CREATE TABLE IF NOT EXISTS dim_products (
  sku VARCHAR(50) NOT NULL,
  product_name VARCHAR(255),
  category VARCHAR(100)
)
DISTSTYLE ALL      -- Use DISTSTYLE ALL because it's a small table that joins to large fact tables.
SORTKEY(sku);

CREATE TABLE IF NOT EXISTS fact_daily_sales (
  date_key DATE,
  sku VARCHAR(50),
  product_name VARCHAR(255),
  total_quantity_sold BIGINT -- Using BIGINT to match the Spark sum() output
)
DISTKEY(sku)
SORTKEY(date_key);

```



Step 8: Create Sns and Subscribe to email

inventory-discrepancy-alerts

EditDeletePublish message

Details

Name

inventory-discrepancy-alerts

ARN

arn:aws:sns:ap-south-1:487615743519:inventory-discrepancy-alerts

Type

Standard

Display name

-

Topic owner

487615743519

Subscriptions

Access policyData protection policyDelivery policy (HTTP/S)Delivery status loggingEncryptionTagsIntegrations

Subscriptions (1)

EditDeleteRequest confirmationConfirm subscriptionCreate subscription

Search

< 1 > ⚙

ID	Endpoint	Status	Protocol
17676cc5-a56a-4d2b-8d85-390ba89b447d	codingmatters2004@gmail.com	Confirmed	EMAIL

AWS Notification - Subscription Confirmation

Inbox x

⌵ ⌵

AWS Notifications

<no-reply@sns.amazonaws.com>
to me

13:23 (7 hours ago) ☆ 😊 ↶ ⋮

You have chosen to subscribe to the topic:
arn:aws:sns:ap-south-1:487615743519:inventory-discrepancy-alerts

To confirm this subscription, click or visit the link below (If this was in error no action is necessary):
[Confirm subscription](#)

Please do not reply directly to this email. If you wish to remove yourself from receiving all future SNS subscription confirmation requests please send an email to [sns-opt-out](#)

Project Resource Summary

Here is a breakdown of the AWS services utilized in this data pipeline and the rationale for their use.

1. IAM (Identity and Access Management)

- **Resource:** An **IAM Role**.
 - **Purpose:** Created to grant the necessary permissions for AWS services to communicate with each other securely. For example, this role allows the AWS Glue jobs to read data from the S3 bucket and write the processed data to the Amazon Redshift warehouse without needing to manage static credentials.
-

2. S3 (Simple Storage Service)

- **Resource:** An **S3 Bucket**.
 - **Purpose:** Serves as the central **data lake** for this project. It is used to store raw source data (like daily sales files) and potentially intermediate, processed data before it's loaded into the data warehouse. Its scalability and cost-effectiveness make it ideal for storing large volumes of data.
-

3. AWS Glue

- **Resource:** Three distinct **AWS Glue ETL jobs**.
- **Purpose:** Used as the serverless Extract, Transform, and Load (ETL) service to process the raw data from S3.
 - **Glue Job (dim_products):** This job is responsible for processing the product-related data to create and populate the dim_products dimension table. It cleans, transforms, and standardizes product information.
 - **Glue Job 1 (fact_daily_sales):** This job processes the daily transactional sales data, aggregating and structuring it to be loaded into the fact_daily_sales table.
 - **Glue Job 2 (reconciliation):** This job performs a reconciliation process, likely comparing inventory or sales data from different sources to ensure accuracy and consistency. The output is loaded into the fact_inventory_reconciliation table.

4. Amazon Redshift Serverless

- **Resource:** A **Redshift Serverless** data warehouse.
- **Purpose:** Acts as the project's central **data warehouse**. It stores the final, structured, and transformed data in a columnar format, which is highly optimized for fast analytical queries and business intelligence (BI) reporting. The serverless option was chosen to automatically scale compute resources based on workload, simplifying management and optimizing costs.

5. Redshift Tables

- **Resource:** Three tables within the Redshift warehouse.
- **Purpose:** To store the cleaned and processed data in a structured, queryable format.
 - **dim_products:** A **dimension table** that stores descriptive attributes about the products (e.g., name, category, brand).
 - **fact_daily_sales:** A **fact table** that stores quantitative sales metrics (e.g., units sold, revenue) on a daily basis.
 - **fact_inventory_reconciliation:** A **fact table** containing the results of the data reconciliation process, highlighting discrepancies or confirming consistency.

6. SNS (Simple Notification Service)

- **Resource:** An **SNS Topic** with an email subscription.
- **Purpose:** Implemented as an alerting and monitoring mechanism. It sends out automated email notifications to subscribers upon the success or failure of the Glue jobs or other pipeline events. This ensures that stakeholders are immediately aware of the pipeline's status.

Apache Airflow?

Airflow is used for **workflow orchestration**. While we have created individual jobs (like your Glue jobs), a real-world project requires you to run them in a specific order, on a schedule, and handle any failures gracefully.

Airflow allows you to:

- **Schedule & Automate:** Automatically run your entire data pipeline at a specific time (e.g., daily at 2 AM) without manual intervention.
- **Manage Dependencies:** Define the exact order of operations. For example, you can ensure your fact_daily_sales and dim_products jobs complete successfully **before** the reconciliation job begins.
- **Monitor & Alert:** Provides a powerful user interface to visualize your pipeline's progress, check logs for each step, and send alerts when tasks fail.
- **Retry & Recover:** Automatically retries a failed task a set number of times, which can resolve temporary issues without waking you up in the middle of the night.
- **Scalability:** Manages complex workflows with many steps and dependencies, which is difficult to do with simple scripts or cron jobs.

Essentially, Airflow acts as the "brain" or "conductor" of your data pipeline, ensuring every part runs correctly, in the right order, and at the right time.

Set Up an Airflow Environment:

Docker Airflow: Build locally, deploy anywhere, and orchestrate with confidence.

Containers [Give feedback](#)

Container CPU usage

19.01% / 400% (4 CPUs available)

Container memory usage

858.5MB / 3.69GB

Show charts

☐ Only show running containers

<input type="checkbox"/>	Name	Container ID	Image	Port(s)	CPU (%)	Memory usage...	Actions
<input type="checkbox"/>	spark	-	-	-	0%	0B / 0B	<input type="play"/> <input type="vertical-ellipsis"/> <input type="trash"/>
<input checked="" type="checkbox"/>	kinesis1-main	-	-	-	19.01%	858.5MB / 15.12G	<input type="play"/> <input type="vertical-ellipsis"/> <input type="trash"/>
<input type="checkbox"/>	airflow-webserver-1	c2dc2ffdc774	apache/airflow:2.7.1	8080:8080	0.27%	325.7MB / 3.78Gi	<input type="play"/> <input type="vertical-ellipsis"/> <input type="trash"/>
<input type="checkbox"/>	airflow-init-1	150c59366423	apache/airflow:2.7.1		0%	0B / 0B	<input type="play"/> <input type="vertical-ellipsis"/> <input type="trash"/>
<input type="checkbox"/>	airflow-scheduler-1	88dc1ba21103	apache/airflow:2.7.1		4.74%	326.7MB / 3.78Gi	<input type="play"/> <input type="vertical-ellipsis"/> <input type="trash"/>
<input type="checkbox"/>	postgres-1	52e1bad2bde7	postgres:13	5432:5432	1.66%	53.3MB / 3.78GB	<input type="play"/> <input type="vertical-ellipsis"/> <input type="trash"/>
<input type="checkbox"/>	mysql-1	49b3dd0fdb7a	mysql:8.0	3306:3306	12.34%	152.8MB / 3.78Gi	<input type="play"/> <input type="vertical-ellipsis"/> <input type="trash"/>

localhost:8080/login/

Gmail

earning

code

https://chat.openai...

YouTube

Python Exercises

Unstop: Connecting...

Outlier

Aws Console

All Bookmarks

Airflow

15:54 UTC

Log In

Sign In

Enter your login and password below:

Username:

Password:

Sign In

Define Connections:

- In the Airflow UI, you configure connections to your external systems. You would create an **AWS Connection** that gives Airflow the permissions (via an IAM AccessKeys) to interact with your S3, Glue, and Redshift services.

List Connection								
Search▼								
+ Actions▼ ← Record Count: 2								
<input type="checkbox"/>	Conn Id	Conn Type	Description	Host	Port	Is Encrypted	Is Extra Encrypted	
<input type="checkbox"/>	aws_default	aws				False	False	
<input type="checkbox"/>	redshift_default	redshift		default-workgroup.487615743519.ap-south-1.redshift-serverless.amazonaws.com	5439	False	False	

Edit Connection	
Connection Id *	aws_default
Connection Type *	Amazon Web Services x
Connection Type missing? Make sure you've installed the corresponding Airflow Provider Package.	
Description	
AWS Access Key ID	AKIAxDCBUYPYNKAHNCZ
AWS Secret Access Key	wJalIXUtnFEMiK7MDENGbPxRfiCYEXAMPLEKEY
Extra	<pre>{ "region_name": "ap-south-1" }</pre>

Edit Connection	
Connection Id *	redshift_default
Connection Type *	Amazon Redshift x
Connection Type missing? Make sure you've installed the corresponding Airflow Provider Package.	
Description	
Host	default-workgroup.487615743519.ap-south-1.redshift-serverless.amazonaws.com
Database	dev
User	admin
Password	
Port	5439
Extra	<pre>{ }</pre>

Airflow Dag:

```
from __future__ import annotations
import pendulum
from airflow.models.dag import DAG
from airflow.operators.empty import EmptyOperator
from airflow.providers.amazon.aws.operators.glue import GlueJobOperator
from airflow.providers.common.sql.operators.sql import SQLExecuteQueryOperator
from airflow.models.param import Param # Import the Param class

# --- Configuration ---
AWS_CONN_ID = "aws_default"
REDSHIFT_CONN_ID = "redshift_default"
IAM_ROLE_ARN = "arn:aws:iam::487615743519:role/RedshiftS3ReadRole"
S3_BUCKET = "my-retail-etl-project"
SNS_TOPIC_ARN = "arn:aws:sns:ap-south-1:487615743519:inventory-discrepancy-alerts"
AWS_REGION = "ap-south-1"

# --- MODIFICATION: Use `params.processing_date` instead of `ds` ---
# This makes the SQL commands use the date provided by the user in the UI.
SQL_DELETE_FACT_DAILY_SALES = "DELETE FROM fact_daily_sales WHERE date_key = '{{\nparams.processing_date\n}}';"
SQL_COPY_FACT_DAILY_SALES = f"""
COPY fact_daily_sales
FROM 's3://{S3_BUCKET}/staging/pos_sales/date={{\nparams.processing_date\n}}/'
IAM_ROLE '{{IAM_ROLE_ARN}}'
FORMAT AS PARQUET;
"""

SQL_TRUNCATE_DIM_PRODUCTS = "TRUNCATE dim_products;"
SQL_COPY_DIM_PRODUCTS = f"""
COPY dim_products
FROM 's3://{S3_BUCKET}/processed/dim_products/'
IAM_ROLE '{{IAM_ROLE_ARN}}'
FORMAT AS PARQUET;
"""

SQL_DELETE_FACT_INVENTORY_RECONCILIATION = "DELETE FROM fact_inventory_reconciliation WHERE date_key\n= '{{\nparams.processing_date\n}}';"
SQL_COPY_FACT_INVENTORY_RECONCILIATION = f"""
COPY fact_inventory_reconciliation
FROM 's3://{S3_BUCKET}/processed/reconciled_inventory/date={{\nparams.processing_date\n}}/'
IAM_ROLE '{{IAM_ROLE_ARN}}'
FORMAT AS PARQUET;
"""

with DAG(
    dag_id="manual_glue_redshift_retail_pipeline",
    start_date=pendulum.datetime(2025, 1, 1, tz="UTC"),
    schedule=None,
    catchup=False,
    doc_md="""
    ## Manual Retail Data Pipeline
    Orchestrates AWS Glue jobs and Redshift loads.
    Click the play button and enter a `processing_date` to run.
    """,
    # --- MODIFICATION: Define a parameter for the user to enter ---
    params={
        "processing_date": Param(
            "2025-10-12", # This is the default value that will appear in the box
            type="string",
```



```

        title="Processing Date",
        description="The date for which to run the pipeline (format: YYYY-MM-DD).",
    ),
    tags=["glue", "redshift", "manual-trigger"],
) as dag:

    start = EmptyOperator(task_id="start")

    # --- MODIFICATION: Use `params.processing_date` in script_args ---
    glue_job_stage_pos_sales = GlueJobOperator(task_id="glue_job_stage_pos_sales", job_name="glue-
job-1-stage-pos-sales", aws_conn_id=AWS_CONN_ID, script_args={"--processing_date": "{{
params.processing_date }}"})
    glue_job_create_dim_products = GlueJobOperator(task_id="glue_job_create_dim_products",
job_name="glue-job-create-dim-products", aws_conn_id=AWS_CONN_ID, script_args={"--processing_date":
"{{ params.processing_date }}"})
    glue_job_reconcile_inventory = GlueJobOperator(
        task_id="glue_job_reconcile_inventory",
        job_name="glue-job-2-reconcile-inventory",
        aws_conn_id=AWS_CONN_ID,
        script_args={
            "--processing_date": "{{ params.processing_date }}",
            "--s3_bucket": S3_BUCKET,
            "--aws_region": AWS_REGION,
            "--sns_topic_arn": SNS_TOPIC_ARN
        }
    )

    # Redshift tasks now use the single-statement SQL variables
    delete_fact_daily_sales = SQLExecuteQueryOperator(task_id="delete_fact_daily_sales",
conn_id=REDSHIFT_CONN_ID, sql=SQL_DELETE_FACT_DAILY_SALES)
    load_fact_daily_sales = SQLExecuteQueryOperator(task_id="load_fact_daily_sales_to_redshift",
conn_id=REDSHIFT_CONN_ID, sql=SQL_COPY_FACT_DAILY_SALES)

    truncate_dim_products = SQLExecuteQueryOperator(task_id="truncate_dim_products",
conn_id=REDSHIFT_CONN_ID, sql=SQL_TRUNCATE_DIM_PRODUCTS)
    load_dim_products = SQLExecuteQueryOperator(task_id="load_dim_products_to_redshift",
conn_id=REDSHIFT_CONN_ID, sql=SQL_COPY_DIM_PRODUCTS)

    delete_fact_inventory_reconciliation =
SQLExecuteQueryOperator(task_id="delete_fact_inventory_reconciliation", conn_id=REDSHIFT_CONN_ID,
sql=SQL_DELETE_FACT_INVENTORY_RECONCILIATION)
    load_fact_inventory_reconciliation =
SQLExecuteQueryOperator(task_id="load_fact_inventory_reconciliation_to_redshift",
conn_id=REDSHIFT_CONN_ID, sql=SQL_COPY_FACT_INVENTORY_RECONCILIATION)

    end = EmptyOperator(task_id="end")

    # --- Dependencies (Unchanged) ---
    start >> [glue_job_stage_pos_sales, glue_job_create_dim_products]
    glue_job_stage_pos_sales >> delete_fact_daily_sales >> load_fact_daily_sales
    glue_job_create_dim_products >> truncate_dim_products >> load_dim_products
    [glue_job_stage_pos_sales, glue_job_create_dim_products] >> glue_job_reconcile_inventory
    glue_job_reconcile_inventory >> delete_fact_inventory_reconciliation >>
load_fact_inventory_reconciliation
    [load_fact_daily_sales, load_dim_products, load_fact_inventory_reconciliation] >> end

```


Choose the Date.

The screenshot displays the Databricks Jobs interface for a job named 'manual_glue_redshift_retail_pipeline' on October 13, 2025, at 14:37:25 UTC. The job is in a 'Failed' state, as indicated by the red 'Failed' label at the top. The job history table on the left shows a sequence of tasks, with the final task 'end' failing. The job details on the right show a graph of the job's execution flow, including tasks like 'glue_job_create_dim_products', 'glue_job_reconcile_inventory', and 'load_fact_inventory_reconciliation'.

Verify in Glue, S3 and Redshift:

glue-job-1-stage-pos-sales

Last modified on 10/13/2025, 7:50:40 PM

Actions Save Run

ScriptJob detailsRunsData qualitySchedulesVersion Control

Job runs (1/10)

Info

Last updated (UTC)
October 13, 2025 at 14:40:28

View details

Stop job run

Troubleshoot with AI

Table View

Card View

Filter job runs by property

Run status	Retries	Start time (Local)	End time (Local)	Duration	Capacity (DPUs)	Worker type	Glue version
Success	0	10/13/2025 20:08:10	10/13/2025 20:09:31	1 m 12 s	2 DPUs	G.1X	5.0
Success	0	10/13/2025 18:35:51	10/13/2025 18:37:31	1 m 27 s	2 DPUs	G.1X	5.0
Success	0	10/13/2025 18:31:30	10/13/2025 18:32:56	1 m 10 s	2 DPUs	G.1X	5.0
Stopped	0	10/13/2025 18:24:38	10/13/2025 18:26:50	1 m 4 s	2 DPUs	G.1X	5.0
Success	0	10/13/2025 18:06:54	10/13/2025 18:08:04	1 m 1 s	2 DPUs	G.1X	5.0
Success	0	10/13/2025 18:01:08	10/13/2025 18:02:43	1 m 18 s	2 DPUs	G.1X	5.0
Success	0	10/13/2025 17:54:56	10/13/2025 17:56:31	1 m 22 s	2 DPUs	G.1X	5.0

glue-job-create-dim-products

Last modified on 10/13/2025, 8:04:51 PM

Actions Save Run

ScriptJob detailsRunsData qualitySchedulesVersion Control

Job runs (1/9)

Info

Last updated (UTC)
October 13, 2025 at 16:06:15

View details

Stop job run

Troubleshoot with AI

Table View

Card View

Filter job runs by property

Run status	Retries	Start time (Local)	End time (Local)	Duration	Capacity (DPUs)	Worker type	Glue version
Success	0	10/13/2025 20:08:10	10/13/2025 20:09:34	1 m 15 s	2 DPUs	G.1X	5.0
Success	0	10/13/2025 18:35:51	10/13/2025 18:37:05	1 m 5 s	2 DPUs	G.1X	5.0
Success	0	10/13/2025 18:31:30	10/13/2025 18:32:59	1 m 2 s	2 DPUs	G.1X	5.0
Success	0	10/13/2025 18:24:38	10/13/2025 18:25:47	1 m 2 s	2 DPUs	G.1X	5.0
Success	0	10/13/2025 18:06:54	10/13/2025 18:08:07	1 m 5 s	2 DPUs	G.1X	5.0
Success	0	10/13/2025 18:01:08	10/13/2025 18:02:44	1 m 19 s	2 DPUs	G.1X	5.0
Success	0	10/13/2025 17:54:56	10/13/2025 17:56:05	1 m 2 s	2 DPUs	G.1X	5.0

glue-job-2-reconcile-inventory

Last modified on 10/13/2025, 8:06:36 PM

Actions Save Run

ScriptJob detailsRunsData qualitySchedulesVersion Control

Job runs (1/16)

Info

Last updated (UTC)
October 13, 2025 at 16:06:39

View details

Stop job run

Troubleshoot with AI

Table View

Card View

Filter job runs by property

Run status	Retries	Start time (Local)	End time (Local)	Duration	Capacity (DPUs)	Worker type	Glue version
Success	0	10/13/2025 20:09:56	10/13/2025 20:11:04	1 m 4 s	2 DPUs	G.1X	5.0
Failed	0	10/13/2025 18:37:44	10/13/2025 18:38:37	49 s	2 DPUs	G.1X	5.0
Failed	0	10/13/2025 18:33:10	10/13/2025 18:34:04	50 s	2 DPUs	G.1X	5.0
Failed	0	10/13/2025 18:26:57	10/13/2025 18:28:15	1 m 6 s	2 DPUs	G.1X	5.0
Failed	0	10/13/2025 18:08:17	10/13/2025 18:09:09	48 s	2 DPUs	G.1X	5.0
Failed	0	10/13/2025 18:03:04	10/13/2025 18:04:25	1 m 11 s	2 DPUs	G.1X	5.0
Failed	0	10/13/2025 17:56:45	10/13/2025 17:58:00	1 m 5 s	2 DPUs	G.1X	5.0

Amazon S3

staging/

Objects (2)

Copy S3 URL

Copy URL

Download

Open

Delete

Actions

Create folder

Upload

Find objects by prefix

Name	Type	Last modified	Size	Storage class
pos_sales_#folder#	-	October 13, 2025, 15:43:58 (UTC+05:30)	-	GB Standard
pos_sales/	Folder	-	-	-

Amazon S3

>

Buckets

>

my-retail-etl-project

>

staging/

>

pos_sales/

Copy S3 URI

pos_sales/

Objects

Properties

Objects (1)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Find objects by prefix

Name

Type

Last modified

Size

Storage class

date=2025-10-13/

Folder

-

-

-

Amazon S3

>

Buckets

>

my-retail-etl-project

>

staging/

>

pos_sales/

>

date=2025-10-13/

Copy S3 URI

date=2025-10-13/

Objects

Properties

Objects (1)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Find objects by prefix

Name

Type

Last modified

Size

Storage class

part-000000-4421d9b6-882c-43d3-b400-faaf6170272a-c000-snappy.parquet

parquet

October 13, 2025, 20:09:19 (UTC+05:30)

1.3 kB

Standard

Amazon S3

>

Buckets

>

my-retail-etl-project

>

processed/

Copy S3 URI

processed/

Objects

Properties

Objects (2)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Find objects by prefix

Name

Type

Last modified

Size

Storage class

dim_products/

Folder

-

-

-

reconciled_inventory/

Folder

-

-

-

Amazon S3

>

Buckets

>

my-retail-etl-project

>

processed/

>

dim_products/

Copy S3 URI

dim_products/

Objects

Properties

Objects (1)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Find objects by prefix

Name

Type

Last modified

Size

Storage class

part-000000-1a146c25-0e47-47ab-b8ca-781a7b06099e-c000-snappy.parquet

parquet

October 13, 2025, 20:09:23 (UTC+05:30)

1.9 kB

Standard

Amazon S3

>

Buckets

>

my-retail-etl-project

>

processed/

>

reconciled_inventory/

Copy S3 URI

reconciled_inventory/

Objects

Properties

Objects (1)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Find objects by prefix

Name

Type

Last modified

Size

Storage class

date=2025-10-13/

Folder

-

-

-

Amazon S3

>

Buckets

>

my-retail-etl-project

>

processed/

>

reconciled_inventory/

>

date=2025-10-13/

Copy S3 URI

date=2025-10-13/

Objects

Properties

Objects (1)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Find objects by prefix

Name

Type

Last modified

Size

Storage class

part-000000-cab16d1e-ad14-4e4d-b093-532a293a3995-c000-snappy.parquet

parquet

October 13, 2025, 20:10:48 (UTC+05:30)

5.9 kB

Standard

Redshift query editor v2

Filter resources

- Serverless: default-workgroup
 - native databases (2)
 - dev
 - public
 - Tables
 - dev_products
 - fact_daily_sales
 - fact_inventory_reconciliation

3

```
1 SELECT
2 *
3 FROM
4 "dev"."public"."fact_inventory_reconciliation";
5
6
7 SELECT
8 *
9 FROM
10 "dev"."public"."dim_products";
11
12
13 SELECT
14 *
15 FROM
16 "dev"."public"."fact_daily_sales";
```

Result 1 (50)

date key	sku	product name	opening stock	quantity sold	expected closing st...	actual closing stock	discrepancy amount
2025-10-12	SKU007	WebCam	50	2	48	48	0
2025-10-12	SKU004	HDMI Cable	300	19	280	290	0
2025-10-12	SKU006	WiFi Adapter	45	3	42	42	0
2025-10-12	SKU045	Headphones	170	4	166	166	0
2025-10-12	SKU018	Tablet 10inch	20	1	19	19	0
2025-10-12	SKU031	Smartphone Stand	130	7	123	123	0
2025-10-12	SKU006	Laptop Stand	110	6	104	104	0
2025-10-12	SKU029	HDMI Splitter	40	1	39	39	0
2025-10-12	SKU014	Printer Ink	40	5	35	35	0
2025-10-12	SKU011	External HDD 1TB	70	3	67	67	0
2025-10-12	SKU021	HD Webcam	45	3	42	42	0
2025-10-12	SKU015	Laptop Bag	100	4	96	96	0
2025-10-12	SKU050	Router	55	2	53	53	0
2025-10-12	SKU005	USB Hub	90	4	86	86	0

Query ID: 1211057 Elapsed time: 315 ms Total rows: 50

Result 1 (50)

sku	product name	category
SKU001	Wireless Mouse	Accessories
SKU002	Keyboard	Accessories
SKU003	Monitor 24inch	Displays
SKU004	HDMI Cable	Cables
SKU005	USB Hub	Accessories
SKU006	Laptop Stand	Accessories
SKU007	WebCam	Electronics
SKU008	Headphones	Audio
SKU009	Mouse Pad	Accessories
SKU010	Wireless Keyboard	Accessories
SKU011	External HDD 1TB	Storage
SKU012	Pen Drive 32GB	Storage
SKU013	Router	Networking
SKU014	Printer Ink	Printers

RunLimit 100ExplainIsolated sessionServerless: de...devSchedule

```
1 SELECT
2 *
3 FROM
4   "dev"."public"."fact_inventory_reconciliation";
5
6
7 SELECT
8 *
9 FROM
10  "dev"."public"."dim_products";
11
12
13 SELECT
14 *
15 FROM
16  "dev"."public"."fact_daily_sales";
```


Result 1 (50)ExportChart

	date key	sku	product name	total quantity sold
<input type="checkbox"/>	2025-10-12	SKU004	HDMI CABLE	10
<input type="checkbox"/>	2025-10-12	SKU036	WIFI ADAPTER	3
<input type="checkbox"/>	2025-10-12	SKU045	HEADPHONES	4
<input type="checkbox"/>	2025-10-12	SKU039	KEYBOARD	5
<input type="checkbox"/>	2025-10-12	SKU048	EXTERNAL HDD 1TB	2
<input type="checkbox"/>	2025-10-12	SKU001	WIRELESS MOUSE	5
<input type="checkbox"/>	2025-10-12	SKU021	HD WEBCAM	3
<input type="checkbox"/>	2025-10-12	SKU005	USB HUB	4
<input type="checkbox"/>	2025-10-12	SKU050	ROUTER	2
<input type="checkbox"/>	2025-10-12	SKU018	TABLET 10INCH	1
<input type="checkbox"/>	2025-10-12	SKU006	LAPTOP STAND	6
<input type="checkbox"/>	2025-10-12	SKU007	WEBCAM	2
<input type="checkbox"/>	2025-10-12	SKU022	MICROPHONE	2
<input type="checkbox"/>	2025-10-12	SKU029	HDMI SPLITTER	1

Query ID 1412358Elapsed time: 29624 msTotal rows: 50

Check Email:

Alert: Inventory Discrepancy Found for 2025-10-13Inbox x



AWS Notifications<no-reply@sns.amazonaws.com>

to me

20:10 (1 hour ago) ☆ 😊 ↩ ⋮

Inventory reconciliation for date 2025-10-13 found discrepancies.

Total items with discrepancies: 4

Example Discrepancies:

- SKU: SKU010, Product: Wireless Keyboard, Discrepancy: 50
- SKU: SKU047, Product: Wireless Keyboard, Discrepancy: -1
- SKU: SKU005, Product: USB Hub, Discrepancy: 50
- SKU: SKU023, Product: Speaker Bluetooth, Discrepancy: 50

Full report available at: s3://my-retail-etl-project/processed/reconciled_inventory/date=2025-10-13/

--

If you wish to stop receiving notifications from this topic, please click or visit the link below to unsubscribe:
<https://sns.ap-south-1.amazonaws.com/unsubscribe.html?SubscriptionArn=arn:aws:sns:ap-south-1:487615743519:inventory-discrepancy-alerts:17676cc5-a56a-4d2b-8d85-390ba89b447d&Endpoint=codingmatters2004@gmail.com>

Please do not reply directly to this email. If you have any questions or comments regarding this email, please contact us at <https://aws.amazon.com/support>

QuickSight:

